

(CS 601.768) Language Model Agents

First (and Last) Homework

For homework deadline and collaboration policy, check the syllabus

Name: Dev Vyas_____

Sources used for your homework, if any: _____

This assignment will assess your understanding of foundational concepts in information retrieval, post-training, reinforcement learning, agent systems, and efficient language-model inference.

How to hand in your written work: via Gradescope.

Collaboration: You may discuss the homework with your classmates, if that helps you understand the problems and various ideas in the class, but you are **not allowed to share problem solutions with any other students. You must write the solutions individually.** The use of AI (e.g., Github CoPilot) is permitted but **must be documented.**

Typesetting: We strongly recommend typesetting your homework, especially if you have sloppy handwriting! :) We will provide a LaTeX template for homework solutions. Please be as succinct as possible in your answers; overly long responses will be penalized.

1 Agent Loops and Inference Economics

1.1 Theory: Maintaining State in an Agent Loop

Suppose an agent repeatedly takes an action, receives new information from its environment, and updates its internal belief state. Consider three ways of constructing the model context at the next step:

- (a) **Rewrite-state:** maintain a compact belief/state summary near the beginning of the prompt and rewrite that summary after every step.
- (b) **Rewrite-suffix:** keep the earlier prefix fixed, but discard reasoning tokens that produced the previous action.
- (c) **Append-only:** never rewrite previous context; append each new observation or state update to the end of the existing trajectory.

Outline one advantage and one disadvantage for each approach. Consider factors such as context footprint, stale or contradictory information, “**context rot**”, human legibility, and inference efficiency where relevant.

Answer:

- (a) *Rewrite-state: The main advantage is that the agent keeps a compact, up-to-date view of what matters, which helps control context growth and reduces stale information. The downside is that rewriting is lossy, so important details can get dropped or distorted, and changing an early part of the prompt can also reduce prompt-cache reuse.*
- (b) *Rewrite-suffix: The advantage is that the agent can throw away old reasoning or tool-use traces that are no longer useful, so the context stays smaller while the stable prefix can still be cached. The downside is that some of that discarded reasoning may contain useful information about why a decision was made, and editing the suffix can also reduce cache continuity from that point onward.*
- (c) *Append-only: The advantage is that it preserves the full trajectory and is usually the most cache-friendly approach because each new step simply extends the previous prompt. The downside is that the context keeps growing and can fill up with outdated, contradictory, or irrelevant information, which increases context rot and can make the model worse at focusing on the current state.*

1.2 Theory: KV Cache and Agent Design

Why does the memory footprint of computing attention grow quadratically with context length? Suppose an API charges cached input tokens at 0.1x the price of uncached input tokens. Which of the three agent designs above does this pricing structure most strongly incentivize, and why? Which design is least compatible with prefix caching?

Answer: In a conventional full-attention prefill, each of L query positions interacts with L key positions, making the attention-score matrix $L \times L$ (per head) and the attention work $O(L^2)$. Materializing that matrix also takes $O(L^2)$ temporary memory; FlashAttention-style kernels can avoid storing it all. The persistent KV cache itself is $O(L)$ in sequence length for fixed model dimensions, not quadratic. That's an important distinction.

*At 0.1x cached-input pricing, **append-only** is most attractive: yesterday's context is literally the prefix of today's, so it can be reused cheaply. **Rewrite-state** is least cache-friendly because modifying a summary near the start invalidates almost everything after it. Rewrite-suffix still preserves the unchanged earlier prefix, so it lands in the middle.*

1.3 Empirical: Dense vs. Mixture-of-Experts Serving

Using vLLM on a single GPU, investigate when a Mixture-of-Experts model provides a serving-throughput advantage over dense models. Use the following three checkpoints:

Checkpoint	Architecture	Role in comparison
allenai/OLMo-2-0425-1B-Instruct	Dense	small dense baseline
allenai/OLMoE-1B-7B-0924-Instruct	MoE	~1B active / 7B total
allenai/OLMo-2-1124-7B	Dense	large dense baseline

Serve all three models using **FP8 weights** and an **FP8 KV cache**. Keep the serving setup otherwise as comparable as possible across models.

Design an experiment that separately characterizes **prefill-dominated** and **decode-dominated** workloads. For prefill jobs, plot on the X axis the length of the input context and on the Y axis the average tokens/second. For decode jobs, plot the number of parallel generations on the X axis. Your experiment should make it possible to understand the effects of model size, number of active parameters, sequence length, and concurrent decoding. Produce two figures that clearly summarize the prefill and decode results, alongside the raw measurements.

Then briefly interpret the results. Explain how the MoE behaves more like the small dense model and how it behaves more like the large dense model. You may reference the OLMoE technical report [Muennighoff et al., 2025] for more information about this model and its performance characteristics.

Include a pointer to the script in your code that implements this experiment.

*Answer: **Setup.** I ran the three specified checkpoints one at a time on a single RTX 4070 SUPER with vLLM 0.30.0. Every model used online per-tensor FP8 weights, an FP8 KV cache, BF16 for unquantized tensors, eager execution, a 4096-token context limit, and disabled prefix caching. For prefill, the same deterministic technical passage was tokenized to 512, 1024, 2048, and 3000 body tokens; a short *yes* instruction followed it, and exactly one output token was forced. I measured prompt tokens divided by end-to-end request time, three repeats per length. For decode, the climate-essay prompt generated exactly 3000 tokens per request at parallelism 1, 2, and 4; I measured aggregate generated tokens divided by batch wall time, twice per setting. Each point in the figures is the sample mean, with sample-standard-deviation bars.*

Model	Prefill at 3000 body tokens	Decode $c = 1$	Decode $c = 4$
Dense 1B	5,929	46.9	206.9
MoE (1B active / 7B total)	28,455	47.7	182.8
Dense 7B	1,576	11.6	32.4

All rates are tokens/second; decode rates sum tokens across requests.




Figure 1: Prefill throughput versus input-context length (model-specific prompt wrappers add a few tokens).




Figure 2: Aggregate decode throughput versus concurrent generations.

Reading the curves. Here the MoE's decode speed is much closer to the small dense model than the 7B dense model: at four parallel requests it reaches 182.8 versus 206.9 and 32.4 tokens/s. That fits the idea that only a subset of experts runs per token. But it still has roughly 7B resident parameters, so it is like the large model in weight-memory pressure; the engine reserved only about 3.11 GiB for its FP8 KV cache, versus 8.15 GiB for dense 1B. Interestingly, its measured prefill is faster than both dense baselines at the longest context. I wouldn't sell that as a universal MoE law: FP8 expert kernels, routing, clock/JIT effects, and the unusually slow 7B long-context point all matter. Three prefill and two decode repeats are a small sample. The controlled inference setup is fair as a serving comparison, but model architecture is not the only moving part.

Files on GitHub: benchmark code; raw measurements; prefill image; decode image.

2 Efficient Post-Training

2.1 Theory: LoRA Parameter and Memory Accounting

Consider Qwen/Qwen2.5-0.5B, which has 24 transformer layers. Suppose we apply LoRA [Hu et al., 2021] with $r = 16$ to all linear transformations in each transformer layer. These consist of the following seven projections:

$$q_{\text{proj}}, k_{\text{proj}}, v_{\text{proj}}, o_{\text{proj}}, \text{gate}_{\text{proj}}, \text{up}_{\text{proj}}, \text{down}_{\text{proj}}.$$

Their dimensions are:

$$\begin{aligned} q_{\text{proj}}, o_{\text{proj}} &: 896 \rightarrow 896, \\ k_{\text{proj}}, v_{\text{proj}} &: 896 \rightarrow 128, \\ \text{gate}_{\text{proj}}, \text{up}_{\text{proj}} &: 896 \rightarrow 4864, \\ \text{down}_{\text{proj}} &: 4864 \rightarrow 896. \end{aligned}$$

For a weight matrix $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$, a rank- r LoRA introduces matrices

$$A \in \mathbb{R}^{r \times d_{\text{in}}}, \quad B \in \mathbb{R}^{d_{\text{out}} \times r}.$$

- How many trainable parameters does LoRA add to one transformer layer? To all 24 layers?
- Assuming trainable parameters and their gradients are both stored in 16-bit precision, how much memory is required to store the LoRA parameters and gradients? Ignore activations and optimizer state.
- For comparison, treat full fine-tuning as updating 0.49 billion parameters. Under the same assumptions, how much memory is required for the full parameters and gradients?
- If AdamW [Loshchilov and Hutter, 2017] additionally maintains two FP32 moment estimates for every trainable parameter, how does the comparison change? State any additional assumptions you make.

You may validate your answer programmatically by constructing the corresponding PEFT LoRA model and inspecting its number of trainable parameters.

Answer: Each projection adds $r(d_{\text{in}} + d_{\text{out}})$ trainable parameters. Adding the seven projections in one layer gives

$$16[2(896 + 896) + 2(896 + 128) + 3(896 + 4864)] = \boxed{366,592}.$$

Across 24 layers, that is $\boxed{8,798,208}$ parameters. At two bytes each for a parameter and its gradient, the LoRA total takes $4(8,798,208) = 35,192,832$ bytes, or $\boxed{33.56 \text{ MiB}}$. Updating 490 million full-model parameters under the same rule takes 1,960,000,000 bytes, or $\boxed{1.825 \text{ GiB}}$.

With two FP32 AdamW moments, add eight bytes per trained parameter, for twelve bytes each in total: LoRA becomes 105,578,496 bytes ($\boxed{100.69 \text{ MiB}}$), while full tuning becomes 5,880,000,000 bytes ($\boxed{5.476 \text{ GiB}}$). Full tuning is about $55.7\times$ larger under either assumption. These are just parameter/gradient/moment storage estimates; I exclude frozen base weights, activations, allocator overhead, possible FP32 master copies, and optimizer metadata. The actual PEFT rank-16 training run independently reported the same 8,798,208 trainable parameters.

Files on GitHub: calculation code; raw text result. No figure is needed for this arithmetic question.

2.2 Empirical: Supervised Fine-Tuning with LoRA

Supervised fine-tune Qwen/Qwen2.5-0.5B using LoRA and full-weight fine-tuning. Apply LoRA to all linear layers in the transformer using PEFT's `target_modules="all-linear"`. Evaluate LoRA ranks in $\{1, 4, 16\}$, and full-weight tuning.

For training data, use the HuggingFaceH4/ultrachat_200k dataset, which contains multi-turn conversations represented as sequences of user and assistant messages. Select a fixed subset of 1000 conversations from `train_sft` for training and use examples from `test_sft` for held-out evaluation.

Keep the training setup as comparable as possible across LoRA ranks including the training examples, number of epochs, batch size, learning rate, and evaluation procedure.

Report:

- the number of trainable parameters for each LoRA rank,

- the peak GPU memory utilization and training time at each rank,
- the training loss and validation loss over the course of training for each rank, and
- a small qualitative comparison of generations from the base model and each fine-tuned model on held-out user prompts.

Your qualitative examples should make it possible to assess whether fine-tuning produced a meaningful change in the model’s ability to follow instructions and respond in a conversational format. The goal is to demonstrate that you can construct an SFT pipeline, apply parameter-efficient fine-tuning to a language model, and compare the effect of LoRA rank; strong benchmark performance is not required.

Include a pointer to the script in your code that implements this experiment.

Answer: I trained Qwen/Qwen2.5-0.5B for one epoch on the saved 1,000-conversation train_sft subset (seed 42), masking non-assistant tokens. LoRA used PEFT target_modules="all-linear" at ranks 1, 4, and 16; the fourth run updated all weights. All runs used BF16, 512-token truncation, batch size 1, gradient accumulation 8, and gradient checkpointing. The LoRA learning rate was $2 \cdot 10^{-4}$ for all ranks; full tuning used $2 \cdot 10^{-5}$. I also sampled 100 fixed conversations from test_sft (seed 42, shuffle buffer 100; no overlap with the training subset), and apply the same assistant-only loss mask and final-512-token truncation for validation.

Run	Trainable parameters	Time (s)	Peak allocated (GiB)	Last train loss
LoRA $r = 1$	549,888	310.40	1.844	1.5673
LoRA $r = 4$	2,199,552	320.34	1.869	1.5520
LoRA $r = 16$	8,798,208	337.18	1.967	1.5471
Full fine-tuning	494,032,768	158.03	4.616	1.5436

The ranks scale exactly as expected. In this implementation the full-weight run was actually faster despite using much more peak memory, so I would not generalize the wall-time ordering to other software/hardware. The loss column above is training loss, and its last logged point is at step 120 (about epoch 0.96). The original run didn’t save intermediate checkpoints, so I evaluated the saved final checkpoints rather than pretending to recover a training-time validation curve. Mean assistant-token cross-entropy on the identical 100 held-out conversations (42,477 target tokens) was: base 1.7732, LoRA $r = 1$ 1.6528, $r = 4$ 1.6375, $r = 16$ 1.6335, and full FT 1.6274 nats/token. All four tuned checkpoints are lower than the base on this sample; the gaps among tuned variants are much smaller than the gap from base. This is a single endpoint comparison on a 100-conversation subset, not a broad benchmark.



Figure 3: Logged training loss across the four one-epoch runs.




Figure 4: Assistant-token loss on the same fixed 100-conversation test_sft subset for the base model and saved final checkpoints.

The fixed subset contains lots of multi-turn follow-ups, numbered practical advice, and passage-based questions. For a small qualitative check, I tried three new prompts: a damaged-tool tracking plan, a follow-up asking for a wheelchair-accessible clothing swap with two hours to set up, and a question based on a repair-cafe notice. All five models can produce structured advice, but several answers are repetitive or miss the two-hour constraint. More importantly, all five get the notice wrong: an electrical lamp cannot be accepted this month, while 11 a.m. needs no booking. These samples do not establish a meaningful improvement in instruction following.

Limit: This is a final-checkpoint validation comparison on a fixed 100-conversation test_sft subset, not the full test split. It does not show validation loss over training; the original training run saved no intermediate checkpoints. The qualitative prompts are separate from this quantitative subset.

Files on GitHub: subset code; training code; generation code; validation code; validation plotting code; exact training subset; fixed held-out test subset; raw summary; raw loss history; training-loss image; validation summary; per-conversation validation results; validation image; side-by-side generations.

3 Language Model Reinforcement Learning

3.1 Theory: Formulating Language-Model Generation as RL

Consider autoregressive generation from a language model.

- Identify the **state**, **action**, **policy**, **transition function**, **reward**, and **trajectory** in this setting.
- Explain why the transition function is usually deterministic even though the policy is stochastic.
- Describe the difference between treating generation as an **MDP** and as a **contextual bandit**. When would each formulation be more natural? How does this distinction relate to outcome versus process rewards?

Answer:

- **State:** the prompt plus all tokens generated so far.
- **Action:** the next token the model chooses.
- **Policy:** the language model's probability distribution over the next token.
- **Transition:** append the chosen token to the current sequence.
- **Reward:** feedback on the generation, either after individual steps or at the end.

- **Trajectory:** the full sequence of states and generated tokens from the prompt to the final response.

The transition is usually deterministic because once a token is chosen, the next state is simply the old sequence with that token appended. The randomness comes from the policy deciding which token to choose, not from what happens after choosing it.

In an MDP, generation is treated as a sequence of decisions, usually one token or reasoning step at a time. This is more natural when intermediate decisions matter and we have process rewards. In a contextual bandit, the prompt is the context and the entire generated response is treated as one action that receives one reward. This is more natural for outcome rewards such as whether the final answer was correct or preferred. So, roughly: process rewards naturally fit the MDP view, while outcome-only rewards make the contextual-bandit view more natural. An MDP can still use only a final outcome reward.

3.2 Theory: REINFORCE and the Policy Gradient

Describe the core training loop of **REINFORCE**: what do we sample, what do we score, and what objective do we optimize? Why is REINFORCE appealingly simple compared with actor-critic methods? What is its central statistical weakness?

Explain why subtracting an action-independent baseline can change the variance of the policy-gradient estimator without changing its expectation.

Answer: REINFORCE samples a full trajectory from the current policy, assigns that trajectory a reward, and updates the model to make actions from high-reward trajectories more likely and actions from low-reward trajectories less likely. In the basic method, the reward directly weights the policy-gradient update.

REINFORCE is appealingly simple compared with actor-critic methods because it does not require a separate critic or value model; it uses the reward directly. Its main weakness is that the gradient estimate can be very noisy, leading to high variance and poor sample efficiency.

A baseline, such as the expected reward, can reduce variance by centering the return: the update depends on whether a trajectory did better or worse than the baseline. If the baseline is independent of the sampled action, subtracting it does not change the expected policy gradient, because its expected contribution cancels.

3.3 Theory: Importance Sampling and Off-Policy Data

Suppose rollouts were generated according to π_{old} , but we now want to optimize a different policy π_{θ} . Explain how importance sampling allows samples from one distribution to estimate an expectation under another distribution. Why can importance ratios become problematic when the two policies differ substantially? Explain the **bias-variance trade-off** introduced by clipping or truncating importance weights.

Answer: Importance sampling lets us reuse rollouts from π_{old} by reweighting each sample based on how likely the new policy π_{θ} would have been to produce it. So samples that are more likely under the new policy count more, and samples that are less likely count less.

The problem is that if the two policies become very different, these ratios can become extremely large or small. A few trajectories can then dominate the update, giving us very high variance. This is especially an issue for LLMs because ratios can accumulate across many generated tokens.

Clipping or truncating the ratios limits how much any one sample can affect training. This lowers variance and makes training more stable, which is the basic idea behind methods like PPO. The trade-off is that once we modify those true ratios, the estimator becomes biased. So clipping basically trades some correctness in expectation for much more stable updates.

3.4 Theory: Proximal Policy Optimization

PPO [Schulman et al., 2017] updates a policy using the importance ratio

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\text{old}}(a_t | s_t)}.$$

What does $r_t > 1$ mean? What does $r_t < 1$ mean? Explain the role of **clipping** in the PPO objective. What kind of update is clipping intended to prevent, and why might an unconstrained policy-gradient update be undesirable even when its estimated advantage has the correct sign?

Answer: If $r_t(\theta) > 1$, the new policy assigns a higher probability to that action than the old policy did. If $r_t(\theta) < 1$, it assigns a lower probability.

Clipping in PPO prevents the policy from benefiting from making that probability change too large. If an action has positive advantage, PPO encourages making it more likely, but only up to a point. If it has negative advantage, PPO encourages making it less likely, again only up to a point.

The goal is to prevent one update from moving the policy too far away from the policy that generated the data. Even if the advantage has the correct sign, its magnitude can be noisy or imperfect, so an unconstrained update could overreact and make the policy collapse or become unstable. PPO trades some aggressiveness for more stable updates.

3.5 Theory: Group Relative Policy Optimization

GRPO [Shao et al., 2024] replaces a learned value-model baseline with a baseline estimated from multiple responses to the same prompt. Suppose G completions for one prompt receive rewards $r_1, \dots, r_G$. Write the standard group-relative advantage used by GRPO.

Then explain:

- Why might a group-relative advantage be preferable to a learned critic?
- What additional cost does GRPO incur? On what types of tasks would GRPO be impossible to use?
- What happens when all responses in a group receive the same reward?

Answer: The standard GRPO advantage is

$$\hat{A}_i = \frac{r_i - \bar{r}}{\sigma_r}, \quad \bar{r} = \frac{1}{G} \sum_{j=1}^G r_j,$$

where σ_r is the standard deviation of the rewards in the group. Each response is judged relative to the other responses generated for the same prompt.

A group-relative advantage can be preferable to a learned critic because it does not require training and maintaining a separate value model. The other sampled responses provide a baseline directly, saving model memory and avoiding errors from a poorly trained critic.

The extra cost is that multiple responses must be generated and scored for every prompt. This can be expensive, especially for long agent trajectories. GRPO is not really usable when the same situation cannot be reset and replayed multiple times, such as a one-shot real-world interaction where each action permanently changes the environment.

If every response receives the same reward, there is no relative signal within the group. After subtracting the mean, all advantages are zero, so the group provides no useful policy update. Implementations also need to handle the zero standard deviation numerically.

3.6 Theory: Modifications to GRPO

Consider three modifications to vanilla GRPO discussed in the reading:

1. **GSPO's sequence-level importance ratios,**
2. **DAPO's dynamic sampling,**
3. **Dr. GRPO's removal of response-length and question-difficulty biases.**

Choose **two** and explain:

- what failure mode of vanilla GRPO motivates the modification,
- how the proposed modification addresses that failure, and
- what new trade-off, if any, it introduces.

Your answer should focus on the mechanism rather than simply describing the algorithm.

Answer: I would choose DAPO dynamic sampling and Dr. GRPO.

***DAPO dynamic sampling:** In vanilla GRPO, if every response to a prompt gets the same reward, all of the relative advantages become zero. Compute is spent generating those responses, but they provide little or no learning signal. DAPO addresses this by dynamically resampling and retaining groups with reward variation, where at least some responses succeed and some fail. This makes training more efficient, but the trade-off is additional sampling cost and a bias toward prompts where the model has a mix of successes and failures. Very easy or extremely hard prompts may be filtered out more often.*

***Dr. GRPO:** Vanilla GRPO normalizes by response length and by the reward variation within each question's group. This can create biases, such as penalizing long wrong answers less strongly and weighting questions differently depending on their*

reward variance. Dr. GRPO removes those normalization terms so updates more directly reflect whether a response did better or worse than the group average.

The trade-off is that removing length normalization means longer responses can contribute more total gradient simply because they contain more tokens. Recent work argues that this is a genuine trade-off: one can remove GRPO's original length bias, but cannot simultaneously make the estimator completely unbiased and completely independent of response length.

3.7 Theory: On-policy vs. Off-policy RL

What is the distinction between **on-policy** and **off-policy** post-training?

Why might one prefer off-policy learning when high-reward trajectories are rare under the current policy? Conversely, why might training on fresh samples from the current policy become valuable as the policy changes during training?

Explain how microbatches in PPO algorithms and stale rollouts in asynchronous RL systems blur the distinction between completely on-policy and off-policy learning.

Answer: On-policy post-training learns from trajectories generated by the current policy, while off-policy training uses trajectories generated by a different or older policy.

Off-policy learning is useful when good trajectories are very rare under the current model. Instead of waiting for the model to discover them, we can train on successful trajectories produced earlier, by another policy, or through search. The downside is distribution mismatch, because those trajectories may not reflect what the current model would actually do.

Fresh on-policy samples become valuable as the model improves because they show the states and mistakes the current model is actually encountering. Old data can become less relevant as the policy changes.

In practice, the distinction is blurry. PPO usually collects a rollout batch and then performs several minibatch updates on it. After the first few updates, that data was generated by an older version of the policy. Similarly, asynchronous RL systems may generate rollouts using slightly stale copies of the model while the learner keeps updating. Thus, many systems described as on-policy are approximately on-policy, with some controlled amount of policy lag.

3.8 Theory: RL for Agents

The RL algorithms above mostly describe a single prompt followed by a single generated completion. Explain what changes when the model instead interacts with an environment over many turns.

In particular, discuss:

- why tool/environment outputs should generally be masked from the policy-gradient loss,
- why variable wall-clock trajectory lengths create an infrastructure problem that does not arise as strongly in ordinary single-turn RL, and

Answer: With an agent, the model is no longer generating one completion and stopping. It generates an action, gets an observation from a tool or environment, then continues for many turns. So the trajectory now mixes model-generated tokens with environment-generated tokens.

Tool/environment outputs should generally be masked from the policy-gradient loss because they are observations or results, not actions chosen by the model. The model-generated reasoning and tool calls around those observations are still trained, while the observations remain in context to influence later decisions.

The infrastructure also gets harder because trajectories can take very different amounts of time. One agent might finish after one quick tool call, while another waits on several searches or a slow code execution. If training waits for the whole batch to finish, GPUs can sit idle waiting for the slowest trajectory. This is why agent-RL systems increasingly implement asynchronous rollouts and load balancing, so completed trajectories can keep moving without being blocked by long-running ones.

4 Information Extraction and Retrieval

4.1 Theory: Sparse vs. Dense Retrieval

What feature(s) of a query would lead you to expect a **sparse** retrieval system (e.g., BM25) to outperform a **dense** retrieval system (e.g., nearest-neighbor search over dense embeddings with FAISS)? What features would lead you to expect better performance from a dense retriever? Give at least one example query of each type and explain your reasoning.

When might a tool like `grep` be a better retrieval mechanism than BM25 or a dense embedding index? When would it clearly be insufficient?

Answer: A sparse retriever like BM25 is usually better when the exact words matter, especially for rare terms, names, error codes, function names, or IDs. For example, a query like “CVE-2024-3094” should favor BM25 because the exact token is much more useful than semantic similarity. Dense retrieval is better when the query is about meaning and the relevant document may use different wording. For example, “ways to reduce model hallucinations” could retrieve documents discussing “factuality” or “grounding” even if they never use the word “hallucination.”

grep can be even better when I already know the exact string, symbol, or pattern I am looking for, especially in a codebase. If I want every use of `process_payment()` or a specific error message, grep is direct, exact, and does not require maintaining an index.

It becomes insufficient when I do not know the exact wording. For a query like “where is authentication handled?”, the relevant code might contain `validate_token` or `session_manager` without ever saying “authentication.” In that case, semantic or indexed retrieval is much more useful.

4.2 Theory: Bi-Encoders vs. Cross-Encoders

What is the architectural difference between a **bi-encoder** and a **cross-encoder**? In a two-stage retrieval system consisting of *retrieval* followed by *reranking*, which architecture is better suited for each stage? Why?

Answer: A bi-encoder encodes the query and each document separately into embeddings, then compares those embeddings. Because document embeddings can be precomputed and indexed, it is very fast and works well for retrieving candidates from a huge corpus.

A cross-encoder takes the query and document together as one input and lets them directly attend to each other before producing a relevance score. That usually gives better relevance judgments, but it is much slower because the model must run separately for every query–document pair.

In a two-stage system:

- **Retrieval:** use a bi-encoder, because it scales efficiently to millions of documents.
- **Reranking:** use a cross-encoder, because after retrieving only the top 50 or 100 candidates, the extra compute can be spent scoring them more accurately.

4.3 Theory: MapReduce over a Corpus

This section considers questions asked about a corpus that is too large to fit into a language model’s context window: all the emails exchanged within a fictional company OpenBrain. For each question, describe how you could answer the question using a **MapReduce**-style procedure.

1. What proportion of the emails within OpenBrain were sent to CEO Stan Allman?
2. What proportion of the emails within OpenBrain discussed the cybersecurity capabilities of their most powerful model, Capricorn-6?
3. Summarize all the emails Stan Allman received related to Capricorn-6’s cybersecurity capabilities.

In your answer, specify the logic of each map call, including its input and output specifications, alongside the logic of the reduce call(s) required. Consider that a language model query could be used for any of these calls, but will incur more computational cost than programmatic logic applied to a document’s metadata.

Answer:

1. Proportion of emails sent to Stan Allman

Map: For each email, inspect only the metadata and emit (`total`, 1). If Stan is a recipient, also emit (`to_stan`, 1). No LLM is needed here.

Reduce: Sum the values for `total` and `to_stan`, then compute `to_stan/total`.

2. Proportion discussing Capricorn-6 cybersecurity

Map: Give each email body to an LLM and ask whether it meaningfully discusses Capricorn-6’s cybersecurity capabilities. Emit (`total`, 1) for every email and (`relevant`, 1) when the answer is yes.

Reduce: Sum `relevant` and `total`, then compute `relevant/total`.

3. Summary of emails Stan received about Capricorn-6 cybersecurity

Map: First filter emails to Stan using metadata. Then give each of those emails to an LLM. If it is relevant to Capricorn-6 cybersecurity, emit (`capricorn_cyber`, short summary).

Reduce: Group all summaries under that key and combine them into one overall summary, removing repetition and preserving important details. If there are too many summaries to fit in context, perform the reduce hierarchically in multiple rounds.

4.4 Theory: Comparing Retrieval Metrics

Consider **Recall@k** and **Mean Reciprocal Rank (MRR)**.

- Write the formula for each metric and explain in words what it measures.
- What are the best- and worst-case values for each metric?
- How do the two metrics differ in their sensitivity to the rank at which a relevant document appears?
- Give an example where two retrieval systems could have the same Recall@k but different MRR. What does this tell you about the information captured by each metric?

Answer: **Recall@k:**

$$\text{Recall@k} = \frac{\text{number of relevant documents in the top } k}{\text{number of relevant documents overall}}.$$

It measures how much of the relevant information the retriever finds within its top k results. Its best value is 1 and worst is 0.
Mean Reciprocal Rank (MRR):

$$\text{MRR} = \frac{1}{N} \sum_{i=1}^N \frac{1}{\text{rank of the first relevant result for query } i}.$$

It measures how quickly the user sees the first relevant result, averaged across queries. Its best value is 1, when a relevant result is always ranked first, and worst is 0, when no relevant result is retrieved.

The main difference is that Recall@k mostly cares whether relevant documents appear somewhere within the top k, while MRR is very sensitive to where the first relevant result appears.

For example, suppose each query has one relevant document and $k = 5$. System A ranks it at position 1, while System B ranks it at position 5. Both have $\text{Recall@5} = 1$, but their reciprocal ranks are 1 and 1/5. So Recall tells us whether we successfully retrieved the relevant information, while MRR also tells us how well we ranked the first useful result.

References

- Niklas Muennighoff, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Jacob Morrison, Sewon Min, Weijia Shi, Pete Walsh, Oyvind Tafjord, Nathan Lambert, et al. Olmo: Open mixture-of-experts language models. In *International Conference on Learning Representations*, volume 2025, pages 62061–62121, 2025.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.